

CS B551 - Assignment 2: Games

Spring 2021

Due: Tuesday March 30, 11:59:59PM Eastern (New York) time

This assignment will give you practice with game playing and probability. Please read the instructions below carefully; we cannot accept any submissions that do not follow the instructions given here. Most importantly: please **start early**, and ask questions on Piazza or in office hours.

Guidelines for this assignment

Coding requirements. For fairness and efficiency, we use a semi-automatic program to grade your submissions. This means you must write your code carefully so that our program can run your code and understand its output properly. In particular:

1. You must code this assignment in Python 3, not Python 2.
2. Make sure to use the program file name we specify.
3. Use the skeleton code we provide, and follow the instructions in the skeleton code (e.g., to not change the parameters of some functions).
4. You may import standard Python modules for routines not related to AI, such as basic sorting algorithms and data structures like queues, as long as they are already installed on the SICE Linux servers.
5. **IMPORTANT:** In addition to testing your programs on your own, test your code on one of the SICE Linux systems such as burrow.sice.indiana.edu using test scripts that we will provide. We will post more details on the test scripts on Q&A Community.

For each of these problems, you will face some design decisions along the way. Your primary goal is to write clear code that finds the correct solution in a reasonable amount of time. To do this, you should give careful thought to the search abstractions, data structures, algorithms, heuristic functions, etc. To encourage innovation, we will conduct a competition to see which program can solve the hardest problems in the shortest time, and a small amount of your grade (or extra credit) may be based on your program's performance in this competition.

Groups. You'll work in a group of 1-3 people for this assignment; we've already assigned you to a group (see details below) according to your preferences. You should only submit **one** copy of the assignment for your team, through GitHub, as described below. All the people on the team will receive the same grade, except in unusual circumstances; we will collect feedback about how well your team functioned in order to detect these circumstances. The requirements for the assignment are the same no matter how many teammates you have, but we expect that teams with more people will submit answers that are significantly more "polished" — e.g., better documented code, faster running times, more thorough answers to questions, etc.

Coding style and documentation. We will not explicitly grade based on coding style, but it's important that you write your code in a way that we can easily understand it. Please use descriptive variable and function names, and use comments when needed to help us understand code that is not obvious.

Report. Please put a report describing your assignment in the Readme.md file in your Github repository. For each problem, please include: (1) a description of how you formulated each problem; (2) a brief description of how your program works; (3) and discussion of any problems you faced, any assumptions, simplifications, and/or design decisions you made. These comments are especially important if your code does not work as well as you would like, since it is a chance to document how much energy and thought you put into your solution.

Academic integrity. We take academic integrity very seriously. To maintain fairness to all students in the class and integrity of our grading system, we will prosecute any academic integrity violations that we discover. *Before beginning this assignment, make sure you are familiar with the Academic Integrity policy of the course, as stated in the Syllabus, and ask us about any doubts or questions you may have.* To briefly summarize, you may discuss the assignment with other people at a high level, e.g. discussing general strategies to solve the problem, talking about Python syntax and features, etc. You may also consult printed and/or online references, including books, tutorials, etc., but you must cite these materials (e.g. in source code comments). We expect that you'll write your own code and not copy anything from anyone else, including online resources. *However, if you do copy something (e.g., a small bit of code that you think is particularly clever), you have to make it explicitly clear which parts were copied and which parts were your own. You can do this by putting a very detailed comment in your code, marking the line above which the copying began, and the line below which the copying ended, and a reference to the source.* Any code that is not marked in this way must be your own, which you personally designed and wrote. You may not share written answers or code with any other students, nor may you possess code written by another student, either in whole or in part, regardless of format.

Part 0: Getting started

For this project, we are assigning you to a team. We will let you change these teams in future assignments. You can find your assigned teammate(s) by logging into IU Github, at <http://github.iu.edu/>. In the upper left hand corner of the screen, you should see a pull-down menu. Select `cs-b551-sp2021`. Then in the box below, you should see a repository called `userid1-a2`, `userid1-userid2-a2`, or `userid1-userid2-userid3-a2`, where the other user ID(s) correspond to your teammate(s). Now that you know their userid(s), you can write them an email at `userid@iu.edu`.

To get started, clone the github repository:

```
git clone git@github.iu.edu:cs-b551-sp2021/your-repo-name-a2
```

If that doesn't work, instead try:

```
git clone https://github.iu.edu/cs-b551-sp2021/your-repo-name-a2
```

where *your-repo-name* is the one you found on the GitHub website above.

Part 1: Pikachu

Pikachu is a popular childhood game in a certain rural midwestern town that requires just a board consisting of a grid of $n \times n$ (with $n \geq 7$) squares, $2n$ white stones, and $2n$ black stones. Each stone is also called a *Pichu*. Initially the board starts empty, except for two rows of white Pichus on the second and third row of the board, and two rows of black Pichus on rows $n - 2$ and $n - 1$. Two players alternate turns, with White going first. When a Pichu reaches the opposite side of the board (i.e. when a Black Pichu reaches row 1 or a white Pichu reaches row n), then the Pichu is marked with an X and becomes a Pikachu. In any given turn, a player can do one of the following:

- Move a single Pichu of his or her color one square forward, left, or right, if that square is empty.
- Move a single Pichu of his or her color to “jump” over a single piece of the opposite color by moving two squares forward, left, or right, if that square is empty. The jumped piece is removed from the board as soon as it is jumped.
- Move a single Pikachu of his or her color any number of squares forward, left, right, or backwards, to an empty square, as long as all squares in between are also empty.

- Move a single Pikachu of his or her color to “jump” over a single piece of the opposite color and landing any number of squares forward, left, right, or backwards, as long as all of the squares between the Pikachu’s start position and jumped piece are empty and all the squares between the jumped piece and the ending position are empty. The jumped piece is removed as soon as it is jumped.

The winner is the player who captures all of the other player’s pieces first.

Your task is to write a Python program that plays Pikachu well. Your program should accept a command line argument that gives the current state of the board as a string of .’s, w’s, W’s, b’s, and B’s, which indicate which squares have no piece, a white Pichu, a white Pikachu, a black Pichu, and a black Pikachu, respectively, in row-major order. For example, if $n = 7$, then the encoding of the start state of the game would be:

```
.....wwwwwwwww.....bbbbbbbbbbbb.....
```

More precisely, your program will be called with four command line parameters: (1) the value of n , (2) the current player (w or b), (3) the state of the board, encoded as above, and (4) a time limit in seconds. Your program should then decide a recommended single move for the given player with the given current board state, and display the new state of the board after making that move. Displaying multiple lines of output is fine as long as the last line has the recommended board state. The time limit is the amount of time that your program should expect to have to make its decision; our testing code will kill your program at that point, and will use whichever was the last move your program recommended. For example, a sample run of your program might look like:

```
[djcran@macbook]$ python3 ./pikachu.py 7 w .....wwwwwwwww.....bbbbbbbbbbbb..... 10
Hmm, I'd recommend moving the Pichu at row 3 column 3 to row 4 column 3.
New board:
.....wwwwwwwww.www..w....bbbbbbbbbbbb.....
```

The competitions. To make things more interesting, we will provide a way for your program to play against other groups’ programs. We will release this code at least 1 week before the deadline. While the majority of your grade will be on correctness, programming style, quality of answers given in comments, etc., a portion may be based on how well your code performs against others, with particularly well-performing programs eligible for prizes including extra credit points.

Hint: Since our grading program only looks at the last solution that you output, your program can output multiple solutions. In other words, you can choose to completely ignore the time parameter passed to your program and simply output multiple answers until you run out of time and we automatically kill your program.

Note: Your code must conform with the interface standards mentioned above! The last line of the output *must* be the new board in the format given, without any extra characters or empty lines. Also, note that your program cannot assume that the game will be run in sequence from start to end; given a current board position on the command line, your code must find a recommended next best move. Your program can write files to disk to preserve state between runs, but should correctly handle the case when a new board state is presented to your program that is unrelated to the last state it saw.

Part 2: The Game of Sebastian

Sebastian¹ is a one-player game of luck and skill. Each turn has four steps:

1. The player rolls five dice.

¹Yes, Sebastian and Pichu are David’s pet birds and yes this is yet another bird-themed problem.

2. The player inspects the dice and chooses any subset (including none or all) and rolls them.
3. The player inspects the dice again again choose any subset and rolls them.
4. The player must assign the outcome to exactly one category on their score card, depending on which five dice are showing after the third roll.

Here are the categories on the score card:

- **Primis:** The player can add the number of dice that show 1 to his or her score.
- **Secundus:** The player can count the number of dice that show 2, multiply by 2, and add to the score.
- **Tertium:** The player can count the number of dice that show 3, multiply by 3, and add to the score.
- **Quartus:** The player can count the number of dice that show 4, multiply by 4, and add to the score.
- **Quintus:** The player can count the number of dice that show 5, multiply by 5, and add to the score.
- **Sextus:** The player can count the number of dice that show 6, multiply by 6, and add to the score.
- **Company:** If the five dice are either 1, 2, 3, 4, 5 or 2, 3, 4, 5, 6, the player can add 40 points to their score. (Note that for this and all other categories, the order of the dice is not important.)
- **Prattle:** If four of the five dice are either 1, 2, 3, 4, or 2, 3, 4, 5, or 3, 4, 5, 6, the player can add 30 points to their score.
- **Squadron:** If three of the dice show the same number, and the other two dice are also the same, the player can add 25 points to their score.
- **Triplex:** If three of the dice are the same, the player can add up the values of all five dice and add this to their score.
- **Quadrupla:** If four of the dice are the same, the player can add up the values of all five dice and add the sum to their score.
- **Quintuplicatam:** If all five dice are the same, the player can add 50 points to their score.
- **Pandemonium:** The sum of all five dice, no matter what they are.

Players can choose which category to fill at the end of each turn, but **each category may be filled only once per game**. A player can also choose to assign a roll to a category that does not match the requirements, but then a zero is entered into that category.

For example, here's what a few moves of the game might look like, with a single player:

- Player rolls dice and gets 1, 2, 3, 4, 5 on the first roll. They decide not to reroll any dice. Player assigns it to **Company** and gets 40 points. (They could have decided to assign it instead to **Prattle** for 30 points, or to one of the first 5 categories to get a score of 1 for **Primis**, 2 for **Secundus**, 3 for **Tertium**, 4 for **Quartus**, or 5 for **Quintus**, or 15 for **Pandemonium**. Or they could have chosen to get a 0 for any other category).
- Player rolls dice and gets 1, 3, 3, 4, 4. They decide to reroll the 1, and they get a 2. They reroll the 2, and get a 6. So the final dice are 6, 3, 3, 4, 4. They could assign this to ~~**Primis** for 1 point~~, **Tertium** for 6 points, or **Quartus** for 8 points, or 6 for **Sextus**, or 20 for **Pandemonium**, or any other category for 0 points. They choose to assign to **Quartus**, and now have 48 points (i.e., 40 from the first turn plus 8 from this turn).

- Player rolls dice and gets 1, 2, 2, 3, 3. They reroll the 1 and get 4, and reroll the 4 and get 5. So the final dice are 5, 2, 2, 3, 3. Their choices now are: **Quintus** for 5 points, ~~Primis for 1 point~~, **Secundus** for 4 points, **Tertium** for 6 points, or 15 for Pandemonium, or 0 in another category. They choose **Tertium** and now have 54 points.
- Player rolls dice and gets 1, 2, 3, 4, 5 on the first roll. Player can't assign it to **Company** since they've already used this category already, so they count it as a **Prattle** instead, and now have 84 points.
- Player rolls dice and gets 1, 2, 3, 4, 5 after their 3 rolls. The player can't take either **Company** or **Prattle**, or **Tertium** or **Quartus** since they have been used, but they could take **Primis** for 1 point, **Secundus** for 2 points, or **Quintus** for 5 points, or 15 for Pandemonium. However, the player decides it's unlikely they'll ever manage to get a **Quintuplicatam**, whereas it's much more likely they'll be able to get a better score for one of those other categories, so they assign to **Quintuplicatam** and get a score of 0.
- And so on...

After 13 turns, all categories are full, and the game ends. If the player managed to get a score of at least 63 totaled across the first six categories (Primis through Sextus), they get a bonus of 35 points added to their score.

Implement a program that plays Sebastian as well as possible, i.e. getting the highest possible score. Your goal is to achieve as high an average score as possible. As with Assignment 1, a small portion of your grade will be based on how well your program works with respect to the rest of the class. To get you started, we've implemented some skeleton code that should be in your cloned repository. You can run it like,

```
python3 ./sebastian.py
```

The skeleton code implements a very naive automatic player that simply rolls once and assigns the roll to whichever category is next available. You'll want to modify the file called `SebastianAutoPlayer.py`. There are helper programs called `sebastian.py` and `SebastianState.py` to keep track of the scoreboard and to roll the dice. While you should look at `SebastianState.py` to understand how these classes work, and you may want to modify `SebastianState.py` for debugging purposes, your final submission should run **using the `SebastianState.py` file that we supplied without any modifications**, or else we won't be able to grade your submission correctly.

Part 3: Document classification

Many practical problems involve sorting textual objects — documents, emails, sentences, tweets, etc. — into two specific categories — spam vs nonspam, important vs unimportant, acceptable vs inappropriate, etc. Naive Bayes classifiers are often used for such problems. They often use a bag-of-words model, which means that each object is represented as just an unordered “bag” of words, with no information about the grammatical structure or order of words in the document. Suppose there are classes A and B . For a given textual object D consisting of words w_1, w_2, \dots, w_n , a Bayesian classifier evaluates decides that D belongs to A by computing the “odds” and comparing to a threshold,

$$\frac{P(A|w_1, w_2, \dots, w_n)}{P(B|w_1, w_2, \dots, w_n)} > 1,$$

where $P(A|w_1, \dots, w_n)$ is the posterior probability that D is in class A . Using the Naive Bayes assumption, the odds ratio can be factored into $P(A)$, $P(B)$, and terms of the form $P(w_i|A)$ and $P(w_i|B)$. These are the parameters of the Naive Bayes model.

Your job is to write a program that estimates the Naive Bayes parameters from training data (where the correct label is given), and then uses these parameters to classify new text objects.

We've provided skeleton code and a sample training and testing dataset to get you started. You can run it like this:

```
python3 ./classify.py tweets.location.train.txt tweets.location.test.txt
```

Hints: Don't worry, at least at first, about whether the "words" in your model are actually words. Just treat every unique space-delimited token you encounter as a "word," even if it's misspelled, a number, a punctuation mark, etc. It may be helpful to ignore tokens that do not occur more than a handful of times, however.

What to turn in

Turn in the three programs on GitHub (remember to **add**, **commit**, **push**) — we'll grade whatever version you've put there as of 11:59PM on the due date. To make sure that the latest version of your work has been accepted by GitHub, you can log into the github.iu.edu website and browse the code online.